

Chapter 1

The 64 bit x86 C Calling Convention

...

This chapter was derived from a document written by Adam Ferrari and later updated by Alan Batson, Mike Lack, Anita Jones, and Aaron Bloomfield

1.1 What is a Calling Convention?

At the end of the previous chapter, we saw a simple example of a subroutine defined in x86 assembly language. In fact, this subroutine was quite simple – it did not modify any registers except EAX (or RAX) (which was needed to return the result), and it did not call any other subroutines. In practice, such simple function definitions are rarely useful. When more complex subroutines are combined in a single program, a number of complicating issues arise. For example, how are parameters passed to a subroutine? Can subroutines overwrite the values in a register, or does the caller expect the register contents to be preserved? Where should local variables in a subroutine be stored? How should results be returned from functions?

To allow separate programmers to share code and develop libraries for use by many programs, and to simplify the use of subroutines in general, programmers typically adopt a common *calling convention*. The calling convention is simply a set of rules that answers the above questions without ambiguity to simplify the definition and use of subroutines. For example, given a set of calling convention rules, a programmer need not examine the definition of a subroutine to determine how parameters should be passed to that subroutine. Furthermore, given a set of calling convention rules, high-level language compilers can be made to follow the rules, thus allowing hand-coded assembly language routines and high-level language routines to call one another.

In practice, even for a single processor instruction set, many calling conventions are possible. In this class we will examine and use one of the most important conventions: the C language calling convention. Understanding this convention will allow you to write assembly language subroutines that are safely callable from C and C++ code, and will also enable you to call C library functions from your assembly language code.

1.2 The C Calling Convention

The C calling convention is based heavily on the use of the hardware-supported stack. To understand the C calling convention, you should first make sure that you fully understand the push, pop, call, and ret instructions – these will be the basis for most of the rules. In this calling convention, subroutine parameters are passed on the stack. Registers are saved on the stack, and local variables used by subroutines are placed in memory on the stack. In fact, this stack-centric implementation of subroutines is not unique to the C language or the x86 architecture. The vast majority of high-level procedural languages implemented on most processors have used similar calling convention.

The calling convention is broken into two sets of rules. The first set of rules is employed by the caller of the subroutine, and the second set of rules is observed by the writer of the subroutine (the “callee”). It should be emphasized that mistakes in the observance of these rules quickly result in fatal program errors; thus meticulous care should be used when implementing the call convention in your own subroutines.

1.3 The Caller’s Rules

The caller should adhere to the following rules when invoking a subroutine:

1. Before calling a subroutine, the caller should save the contents of certain registers that are designated caller-saved. The caller-saved registers are r10, r11, and any registers that parameters are put into. If you want the contents of these registers to be preserved across the subroutine call, push them onto the stack.
2. To pass parameters to the subroutine, we put up to six of them into registers (in order: rdi, rsi, rdx, rcx, r8, r9). If there are more than six parameters to the subroutine, then push the rest onto the stack in *reverse order* (i.e. last parameter first) – since the stack grows down, the first of the extra parameters (really the seventh parameter) parameter will be stored at the lowest address (this inversion of parameters was historically used to allow functions to be passed a variable number of parameters).
3. To call the subroutine, use the `call` instruction. This instruction places the return address on top of the parameters on the stack, and branches to the subroutine code.
4. After the subroutine returns, (i.e. immediately following the call instruction) the caller must remove any additional parameters (beyond the six stored in registers) from stack. This restores the stack to its state before the call was performed.
5. The caller can expect to find the return value of the subroutine in the register RAX.
6. The caller restores the contents of caller-saved registers (r10, r11, and any in the parameter passing registers) by popping them off of the stack. The caller can assume that no other registers were modified by the subroutine.

Due to the way the calling convention is structured, it will typically be the case that some (or most) of these steps will not make any changes to the stack. For example, if there are six or fewer parameters, then nothing is pushed onto the stack in that step. Likewise, programmers (and compilers) typically keep the results they care about out of the caller-saved registers in steps 1 and 6 to prevent excess pushes and pops.

1.4 The Callee's Rules

The definition of the subroutine should adhere to the following rules:

1. Allocate local variables by using registers or making space on the stack. Recall, the stack grows down, so to make space on the top of the stack, the stack pointer should be decremented. The amount by which the stack pointer is decremented depends on the number of local variables needed. For example, if a local `float` and a local `long` (12 bytes total) were required, the stack pointer would need to be decremented by 12 to make space for these local variables:

Listing 1.1: x86 callee code, part 2

```
sub  rsp , 12
```

As with parameters, local variables will be located at known offsets from the stack pointer.

2. Next, the values of any registers that are designated callee-saved that will be used by the function must be saved. To save registers, push them onto the stack. The callee-saved registers are RBX, RBP, and R12 through R15 (RSP will also be preserved by the call convention, but need not be pushed on the stack during this step).

After these three actions are performed, the actual operation of the subroutine may proceed. When the subroutine is ready to return, the call convention rules continue:

3. When the function is done, the return value for the function should be placed in RAX if it is not already there.
4. The function must restore the old values of any callee-saved registers (RBX, RBP, and R12 through R15) that were modified. The register contents are restored by popping them from the stack. Note, the registers should be popped in the inverse order that they were pushed.
5. Next, we deallocate local variables. The easiest way to do this is to add to RSP the same amount that was subtracted from it in step 1.
6. Finally, we return to the caller by executing a `ret` instruction. This instruction will find and remove the appropriate return address from the stack.

If you look at the assembly generated by some compilers, you will see a few extra commands in there in the callee's prologue:

Listing 1.2: x86 extraneous codedd

```
push  rbp           ; at the start of the callee
mov   rbp , rsp
...
pop   rbp           ; just before the ending 'ret'
```

This code is unnecessary, and is a hold-over from the 32-bit calling convention. You can tell the compiler to not include this code by invoking it with the `-fomit-frame-pointer` flag.

It might be noted that the callee's rules fall cleanly into two halves that are basically mirror images of one another. The first half of the rules apply to the beginning of the function, and are therefor commonly said to define the *prologue* to the function. The latter half of the rules apply to the end of the function, and are thus commonly said to define the *epilogue* of the function.

1.5 Calling Convention Example

The above rules may seem somewhat abstract on first examination. In practice, the rules become simple to use when they are well understood and familiar. To start the process of better understanding the call convention, we now examine a simple example of a subroutine call and a subroutine definition.

Listing 1.3: Example function call, caller's rules obeyed

```
; Want to call a function "myFunc" that takes three
; integer parameters. First parameter is in rax.
; Second parameter is the constant 123. Third
; parameter is in memory location "var"

push rdi                ; rdi will be a param, so saving it

                        ; long retVal = myFunc(x, 123, z);
mov rdi, rax             ; put first param in rdi
mov rsi, 123             ; put second param in rsi
mov rdx, [var]           ; put third param in rdx

call myFunc              ; call the function

pop rdi                  ; restore saved rdi value

; return value of myFunc is now available in rax
; (if there is any return value)
```

In Listing 1.3 a sample function call is depicted. The three parameters are put into the parameter passing registers; if there were more than 6, then the additional ones would be pushed onto the stack in reverse order. The call instruction is used to jump to the beginning of the subroutine in anticipation of the fact that the subroutine will use the ret instruction to return when the subroutine completes. When the subroutine returns, the parameters must be removed from the stack. A simple way to do this is to add the appropriate amount to the stack pointer (since the stack grows down). Finally, the result is available in RAX.

Next up is the caller's rules. An example subroutine implementation that obeys the callee's rules is depicted in Listing 1.4. The subroutine prologue performs the standard actions of allocating local variables by decrementing the stack pointer, and saving register values on the stack.

In the body of the subroutine we can how the local variables are accessed. The prologue put 24 bytes of "stuff" onto the stack: three 8 byte values. The first was the local variable (via the `sub rsp, 8` call). The second was the rbx backup, and the third was the rbp backup. Thus, the stack pointer is now 16 bytes below the local variable, as two 8 byte "things" have been pushed onto the stack since the local variable was allocated. Thus, to access the local variable, one uses `[rsp+16]`, as seen throughout the code.

The function epilogue, as expected, is basically a mirror image of the function prologue. The caller's register values are recovered from the stack, the local variables are deallocated by resetting the stack pointer, and the ret instruction is used to return to the appropriate code location in the caller.

Listing 1.4: Example function definition, callee's rules obeyed

```

global myFunc

section .text

myFunc:
    ; *** Standard subroutine prologue ***
    sub rsp, 8          ; room for a 64-bit local var (result)
    push rbx            ; save callee-save registers
    push rbp            ; both will be used by myFunc

    ; *** Subroutine Body ***
    mov rax, rdi        ; param 1 to rax
    mov rbp, rsi        ; param 2 to rbp
    mov rbx, rdx        ; param 3 to rbx
    mov [rsp+16], rbx   ; put rbx into local var
    add [rsp+16], rbp   ; add rbp into local var
    mov rax, [rsp+16]   ; mov contents of local var to rax
                        ; (return value/final result)

    ; *** Standard subroutine epilogue ***
    pop rbp             ; recover callee save registers
    pop rbx             ; REVERSE of when pushed
    add rsp, 8          ; deallocate local var(s)
    ret                ; pop top value from stack, jump there

```

A good way to visualize the operation of the calling convention is to draw the contents of the nearby region of the stack during subroutine execution. Figure 1.1 depicts the contents of the stack during the execution of the body of `myFunc` (`myFunc` is depicted in Listing 1.4). Notice, lower addresses are depicted lower in the figure, and thus the “top” of the stack is the bottom-most cell. This corresponds visually to the intuitive statement that the x86 hardware stack “grows down.” The cells depicted in the stack are 64-bit wide memory locations, thus the memory addresses of the cells are 4 bytes apart. From this picture we see clearly why the first parameter resides at an offset of 8 bytes from the base pointer. Above the parameters on the stack (and below the base pointer), the call instruction placed the return address, thus leading to an extra 4 bytes of offset from the base pointer to the first parameter.

The assembly code for `myFunc()` was shown above in Listing 1.4. The C++ code to call that subroutine is shown in Listing 1.5.

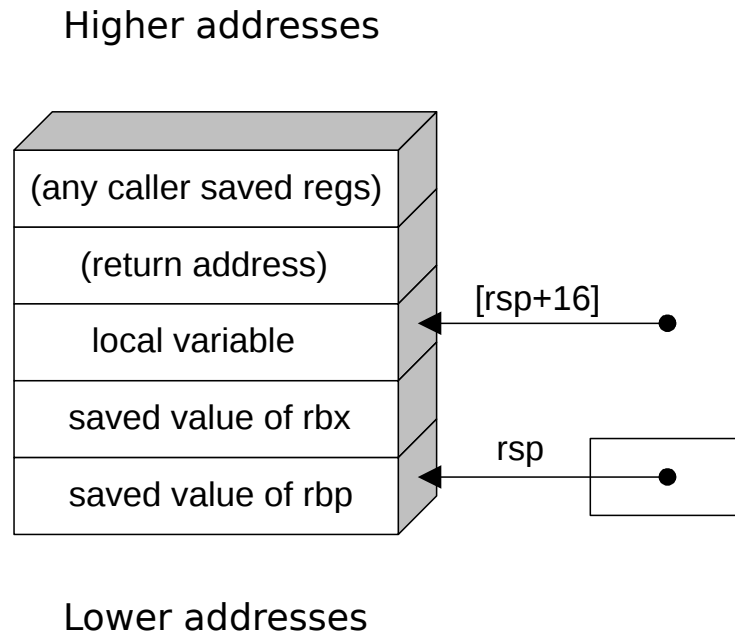


Figure 1.1: A picture of the stack in memory during the execution of the body of `myFunc`

Listing 1.5: Example C++ code to invoke a 3-parameter x86 subroutine

```
#include <iostream>
using namespace std;

extern "C" int myFunc(int, int, int);

int main() {
    int x = 3;
    cout << "myFunc() returned: _"
         << myFunc(x, 5, 10) << endl;
    return 0;
}
```